

# MySQL学习笔记 ( Day031 : 锁\_4 )

MySQL学习

- MySQL学习笔记 ( Day031 : 锁\_4 )
  - 一. 插入意向锁 ( insert intention lock )
    - 1.1. 插入的过程
    - 1.2. 插入意向锁的演示
    - 1.3. 插入意向锁提高插入的并发性演示
  - 二. Read Committed
    - 2.1. 非索引列的等值查询
    - 2.2. 二级索引列的等值查询
  - 三. 锁的算法 ( 二 )
  - 四. 一致性的非锁定读
  - 五. 死锁
    - 5.1. 死锁的介绍
    - 5.2. AB-BA死锁演示

## 一. 插入意向锁 ( insert intention lock )

- 插入意向锁 本质上就是个 Gap Lock
  - 普通Gap Lock 不允许 在 ( 上一条记录: 本记录 ) 范围内插入数据
  - 插入意向锁Gap Lock 允许 在 ( 上一条记录: 本记录 ) 范围内插入数据
- 插入意向锁的作用是为了 提高并发插入的性能 , 多个事务 同时写入 不同数据 至同一索引范围 ( 区间 ) 内 , 并不需要等待其他事务完成 , 不会发生锁等待

### 1.1. 插入的过程

假设现在有记录 10. 30. 50. 70 ; 且为 主键 , 需要插入记录 25 .

- 找到 小于等于25的记录 , 这里是 10
- 找到 记录10的下一条记录 , 这里是 30
- 判断 下一条记录30 上是否有锁 ( 如果有=25的情况 , 后面再讨论 )

- 判断 30 上面如果没有锁 , 则 可以插入
  - 判断 30 上面如果有 Record Lock , 则 可以插入
  - 判断 30 上面如果有 Gap Lock / Next-Key Lock , 则无法插入 , 因为锁的范围是 ( 10, 30 ) / ( 10, 30 ] ; 在 30 上增加 insert intention lock ( 此时处于 waiting 状态 ) , 当 Gap Lock / Next-Key Lock 释放时 , 等待的事物 ( transaction ) 将被 唤醒 , 此时 记录30 上才能获得 insert intention lock , 然后再插入 记录25
- 注意 : 一个事物 insert 25 且没有提交 , 另一个事物 delete 25 时 , 记录25上会有 Record Lock

### 1.2. 插入意向锁的演示

```
--
-- 终端会话1
--
mysql> set tx_isolation="REPEATABLE-READ";
Query OK, 0 rows affected (0.00 sec)

mysql> desc t_lock_1;
+-----+
| Field | Type | Null | Key | Default | Extra |
+-----+
| a | int(11) | NO | PRI | NULL | |
+-----+
1 row in set (0.00 sec)

mysql> begin;
Query OK, 0 rows affected (0.00 sec)

mysql> select * from t_lock_1 where a<=13 for update; -- RR默认认为Next-Key Lock. status与之前一致
+-----+
| a |
+-----+
| 10 |
| 11 |
| 13 |
+-----+
3 rows in set (0.00 sec)

--
-- 终端会话2
--
mysql> set tx_isolation="REPEATABLE-READ";
Query OK, 0 rows affected (0.00 sec)

mysql> begin;
Query OK, 0 rows affected (0.00 sec)

mysql> set innodb_lock_wait_timeout=60; -- 将锁等待时间改成60秒, 原来配置中为5秒, 来不及切换测试
Query OK, 0 rows affected (0.00 sec)

mysql> insert into t_lock_1 values (12);
-- waiting..... (被阻塞了, 在这里等待)

--
-- 终端会话3
--
mysql> show engine innodb status\G
-- -----引擎部分输出-----
---TRANSACTION 25645, ACTIVE 4 sec inserting
mysql tables in use 1, locked 1
LOCK WAIT 2 lock struct(s), heap size 1136, 1 row lock(s)
MySQL thread id 2, OS thread handle 140512469145344, query id 161 localhost root update
-----等待插入的SQL语句-----
insert into t_lock_1 values (12) -- 等待插入的SQL语句
-----
----- TRX HAS BEEN WAITING 4 SEC FOR THIS LOCK TO BE GRANTED:
RECORD LOCKS space id 140 page no 3 n bits 72 index PRIMARY of table 'burn_test'.'t_lock_1' trx id 25645 lock_mode X locks gap before rec insert intention waiting -- 插入记录12的事物等待中 (被终端会话1中的事物阻塞了) , 等待 获得 插入意向锁
Record lock, heap no 4 PHYSICAL RECORD: n_fields 3; compact format; info bits 0
0: len 4; hex 8000000d; asc ;;
1: len 6; hex 000000000602b; asc '+';;
2: len 7; hex ab0000000470128; asc G ;;

-----
TABLE LOCK table 'burn_test'.'t_lock_1' trx id 25645 lock mode IX
RECORD LOCKS space id 140 page no 3 n bits 72 index PRIMARY of table 'burn_test'.'t_lock_1' trx id 25644 lock_mode IX
Record lock, heap no 4 PHYSICAL RECORD: n_fields 3; compact format; info bits 0
0: len 4; hex 8000000d; asc ;;
1: len 6; hex 000000000602b; asc '+';;
2: len 7; hex ab0000000470128; asc G ;;

---TRANSACTION 25644, ACTIVE 36 sec -- 以下是 终端会话1中, <=13 for update产生的Next-Key Lock信息
2 lock struct(s), heap size 1136, 4 row lock(s)
MySQL thread id 5, OS thread handle 140512468612864, query id 113 localhost root cleaning up
TABLE LOCK table 'burn_test'.'t_lock_1' trx id 25644 lock_mode IX
RECORD LOCKS space id 140 page no 3 n bits 72 index PRIMARY of table 'burn_test'.'t_lock_1' trx id 25644 lock_mode X
Record lock, heap no 2 PHYSICAL RECORD: n_fields 3; compact format; info bits 0
0: len 4; hex 8000000d; asc ;;
1: len 6; hex 000000000602b; asc '+';;
2: len 7; hex ab0000000470110; asc G ;;

Record lock, heap no 3 PHYSICAL RECORD: n_fields 3; compact format; info bits 0
0: len 4; hex 8000000b; asc ;;
1: len 6; hex 000000000602b; asc '+';;
2: len 7; hex ab000000047011c; asc G ;;

Record lock, heap no 4 PHYSICAL RECORD: n_fields 3; compact format; info bits 0
0: len 4; hex 8000000d; asc ;;
1: len 6; hex 000000000602b; asc '+';;
2: len 7; hex ab0000000470128; asc G ;;

Record lock, heap no 5 PHYSICAL RECORD: n_fields 3; compact format; info bits 0
0: len 4; hex 80000014; asc ;;
1: len 6; hex 000000000602b; asc '+';;
2: len 7; hex ab0000000470134; asc G 4;;

--
-- 终端会话1
--
mysql> commit; -- 终端会话1 中的事物提交, 相当与释放了锁
Query OK, 0 rows affected (0.00 sec)

--
-- 终端会话2
--
mysql> insert into t_lock_1 values(12); -- 这个是刚才输入的SQL语句, 之前是阻塞的
Query OK, 1 row affected (17.40 sec) -- 因为终端会话1中的事物提交了, 所以此时插入成功

--
-- 终端会话3
--
mysql> show engine innodb status\G
-- -----引擎部分输出-----
2 lock struct(s), heap size 1136, 1 row lock(s), undo log entries 1
MySQL thread id 8, OS thread handle 140333220423424, query id 300 localhost root cleaning up
TABLE LOCK table 'burn_test'.'t_lock_1' trx id 26177 lock_mode IX
RECORD LOCKS space id 140 page no 3 n bits 72 index PRIMARY of table 'burn_test'.'t_lock_1' trx id 26177 lock_mode X locks gap before rec insert intention -- 插入意向锁已经获得了
Record lock, heap no 4 PHYSICAL RECORD: n_fields 3; compact format; info bits 0
0: len 4; hex 8000000d; asc ;; -- 在a=13的记录上有插入意向锁
1: len 6; hex 000000000602b; asc '+';;
2: len 7; hex ab0000000470128; asc G ;;
```

### 1.3. 插入意向锁提高插入的并发性演示

```
--
-- 终端会话1
--

mysql> set tx_isolation='REPEATABLE-READ';
Query OK, 0 rows affected (0.00 sec)


mysql> create table t_lock_5 (a int ,primary key(a));
Query OK, 0 rows affected (0.14 sec)


mysql> insert into t_lock_5 values(5),(20),(50);
Query OK, 3 rows affected (0.00 sec)
Records: 3 Duplicates: 0 Warnings: 0


mysql> begin;
Query OK, 0 rows affected (0.04 sec)


mysql> select * from t_lock_5 where a<=50 for update;
+-----+
| a |
+-----+
| 5 |
| 20 |
| 50 |
+-----+
3 rows in set (0.00 sec)


--
-- 终端会话2
--

mysql> begin;
Query OK, 0 rows affected (0.00 sec)


mysql> insert into t_lock_5 values(25);
-- waiting... 被 终端会话1 中的事物阻塞


--
-- 终端会话1
--

mysql> commit;
Query OK, 0 rows affected (0.00 sec)


--
-- 终端会话2
--

mysql> insert into t_lock_5 values(25);
Query OK, 1 row affected (30.46 sec) -- 由于 终端会话1 中的事物提交了，现在获得了插入意向锁


--
-- 终端会话3
--

mysql> show engine innodb status\G
-- -----省略部分输出-----
2 lock struct(s), heap size 1136, 1 row lock(s), undo log entries 1
MySQL thread id 9, OS thread handle 140333220157184, query id 402 localhost root cleaning up
TABLE LOCK table 'burn_test'.'t_lock_5' trx id 26192 lock mode IX
RECORD LOCKS space id 147 page no 3 n bits 72 index PRIMARY of table 'burn_test'.'t_lock_5' trx id 26192 lock_mode X locks gap before rec insert intention -- 插入意向锁已经获得
Record lock, heap no 4 PHYSICAL RECORD: n_fields 3; compact format; info bits 0
0: len 4; hex 80000032; asc 2;; -- 记录a=50上已经有了插入意向锁
1: len 6; hex 00000000664a; asc fJ;;
2: len 7; hex c30000002a0128; asc * (;;


--
-- 终端会话1
--

mysql> begin; -- 在终端会话1中再开启一个新的事物（此时终端会话2中的事物其实还未提交）
Query OK, 0 rows affected (0.00 sec)


mysql> insert into t_lock_5 values(30); -- 插入小于50, 大于20的值, 插入成功
Query OK, 1 row affected (0.00 sec)


mysql> insert into t_lock_5 values(31); -- 插入小于50, 大于20的值, 插入成功
Query OK, 1 row affected (0.00 sec)


mysql> commit;
Query OK, 0 rows affected (0.02 sec)


-- 这样就可以并发的插入记录了，而不需要一个事物等待另一事物
-- 当所有相关的插入的事物都提交后，50上的插入意向锁 便会释放
```

二. Read Committed

```
--
-- 终端会话1
--

mysql> set tx_isolation='READ-COMMITTED';
Query OK, 0 rows affected (0.00 sec)


mysql> select * from t_lock_1 where a <= 13 for update;
+-----+
| a |
+-----+
| 10 |
| 11 |
| 13 |
+-----+
3 rows in set (0.00 sec)


--
-- 终端会话2
--

mysql> set tx_isolation='REPEATABLE-READ';
mysql> show engine innodb status\G
-- -----省略部分输出-----
---TRANSACTION 26212, ACTIVE 26 sec
2 lock struct(s), heap size 1136, 3 row lock(s)
MySQL thread id 16, OS thread handle 140333220423424, query id 625 localhost root cleaning up
TABLE LOCK table 'burn_test'.'t_lock_1' trx id 26212 lock mode IX
RECORD LOCKS space id 140 page no 3 n bits 72 index PRIMARY of table 'burn_test'.'t_lock_1' trx id 26212 lock_mode X locks rec but not gap -- 在RC隔离级别下，变成了 记录锁（同时会出现幻读问题）
Record lock, heap no 2 PHYSICAL RECORD: n_fields 3; compact format; info bits 0
0: len 4; hex 8000000a; asc ;;
1: len 6; hex 00000000602b; asc '+';;
2: len 7; hex ab000000470110; asc G ;;

Record lock, heap no 3 PHYSICAL RECORD: n_fields 3; compact format; info bits 0
0: len 4; hex 8000000b; asc ;;
1: len 6; hex 00000000602b; asc '+';;
2: len 7; hex ab00000047011c; asc G ;;

Record lock, heap no 4 PHYSICAL RECORD: n_fields 3; compact format; info bits 0
0: len 4; hex 8000000d; asc ;;
1: len 6; hex 00000000602b; asc '+';;
2: len 7; hex ab000000470128; asc G ;;
```

在大部分情况下，RC隔离级别下没有GAP锁，但是一些场景下，线上仍可能出现设置了RC隔离级别后，binlog\_format = row，否则会出现主从不一致的情况

2.1. 非索引列的等值查询

```
--
-- 终端会话1
--

mysql> desc t_lock_3;
+-----+
| Field | Type | Null | Key | Default | Extra |
+-----+
| a | int(11) | NO | | NULL | | -- 没有显示的创建主键
+-----+
1 row in set (0.00 sec)


mysql> set tx_isolation='READ-COMMITTED';
Query OK, 0 rows affected (0.00 sec)


mysql> select * from t_lock_3 where a = 13 for update;
+-----+
| a |
+-----+
| 13 |
+-----+
1 row in set (0.00 sec)


--
-- 终端会话2
--

mysql> show engine innodb status\G
-- -----省略部分输出-----
2 lock struct(s), heap size 1136, 1 row lock(s)
MySQL thread id 16, OS thread handle 140333220423424, query id 632 localhost root cleaning up
TABLE LOCK table 'burn_test'.'t_lock_3' trx id 26213 lock mode IX
RECORD LOCKS space id 143 page no 3 n bits 72 index GEN_CLUST_INDEX of table 'burn_test'.'t_lock_3' trx id 26213 lock_mode X locks rec but not gap -- 只有一个记录锁，且锁住的是GEN_CLUST_INDEX
Record lock, heap no 4 PHYSICAL RECORD: n_fields 4; compact format; info bits 0
0: len 6; hex 000000001206; asc ;; -- 主键
1: len 6; hex 000000006250; asc bP;;
2: len 7; hex c50000002a012c; asc * ,;;
3: len 4; hex 8000000d; asc ;; -- a = 13
```

当查询条件是非索引列的等值查询，RR 隔离级别下会锁住每个记录（形成表锁的效果），而 RC 隔离级别下只锁住查询条件的记录本身

2.2. 二级索引列的等值查询



```
--
-- 终端会话1
--

mysql> alter table t_lock_3 add index idx_a (a);
Query OK, 0 rows affected (0.14 sec)
Records: 0  Duplicates: 0  Warnings: 0

mysql> set tx_isolation='READ-COMMITTED';
Query OK, 0 rows affected (0.00 sec)

mysql> begin;
Query OK, 0 rows affected (0.00 sec)

mysql> select * from t_lock_3 where a = 13 for update;
+-----+
| a |
+-----+
| 13 |
+-----+
1 row in set (0.00 sec)

--
-- 终端会话2
--

mysql> set tx_isolation='READ-COMMITTED';
Query OK, 0 rows affected (0.00 sec)

mysql> show engine innodb status\G
-----省略部分输出-----
3 lock struct(s), heap size 1136, 2 row lock(s)
MySQL thread id 15, OS thread handle 140333220423424, query id 641 localhost root cleaning up
TABLE LOCK table 'burn_test'.'t_lock_3' trx id 26222 lock_mode IX
RECORD LOCKS space id 143 page no 4 n bits 72 index idx_a of table 'burn_test'.'t_lock_3' trx id 26222 lock_mode X locks rec but not gap -- index idx_a 二级索引上的 Record Lock
Record lock, heap no 4 PHYSICAL RECORD: n_fields 2; compact format; info bits 0
0: len 4; hex 8000000d; asc  ;;
1: len 6; hex 00000001206; asc  ;;

RECORD LOCKS space id 143 page no 3 n bits 72 index GEN_CLUST_INDEX of table 'burn_test'.'t_lock_3' trx id 26222 lock_mode X locks rec but not gap -- index GEN_CLUST_INDEX 聚集索引上的 Record Lock
Record lock, heap no 4 PHYSICAL RECORD: n_fields 4; compact format; info bits 0
0: len 6; hex 00000001206; asc  ;;
1: len 6; hex 00000006250; asc  bP;;
2: len 7; hex c5000002a012c; asc  * ,;;
3: len 4; hex 8000000d; asc  ;;

RECORD LOCKS space id 143 page no 3 n bits 72 index GEN_CLUST_INDEX of table 'burn_test'.'t_lock_3' trx id 26222 lock_mode X locks rec but not gap -- index GEN_CLUST_INDEX 聚集索引上的 Record Lock
Record lock, heap no 5 PHYSICAL RECORD: n_fields 2; compact format; info bits 0
0: len 4; hex 80000014; asc  ;;
1: len 6; hex 00000001207; asc  ;;
```

当查询条件是二级索引的等值查询时

- 1. RC 模式下，二级索引查询的记录上有一个记录锁，对应的聚集索引上有一个记录锁
- 2. RR 模式下，二级索引查询的记录上有一个 Next-Key Lock，该记录的下一个记录上有一个 Gap-Lock（二级索引）；对应的聚集索引上有一个记录锁

```
--
-- 终端会话1
--

mysql> set tx_isolation='REPEATABLE-READ';
Query OK, 0 rows affected (0.00 sec)

mysql> begin;
Query OK, 0 rows affected (0.00 sec)

mysql> select * from t_lock_3 where a = 13 for update;
+-----+
| a |
+-----+
| 13 |
+-----+
1 row in set (0.00 sec)

--
-- 终端会话2
--

mysql> set tx_isolation='REPEATABLE-READ';
Query OK, 0 rows affected (0.00 sec)

mysql> show engine innodb status\G
-- -----省略部分输出-----
4 lock struct(s), heap size 1136, 3 row lock(s)
MySQL thread id 16, OS thread handle 140333220423424, query id 649 localhost root cleaning up
TABLE LOCK table 'burn_test'.'t_lock_3' trx id 26223 lock_mode IX
RECORD LOCKS space id 143 page no 4 n bits 72 index idx_a of table 'burn_test'.'t_lock_3' trx id 26223 lock_mode X -- index idx_a 二级索引上的 Next-Key Lock
Record lock, heap no 4 PHYSICAL RECORD: n_fields 2; compact format; info bits 0
0: len 4; hex 8000000d; asc  ;;
1: len 6; hex 00000001206; asc  ;;

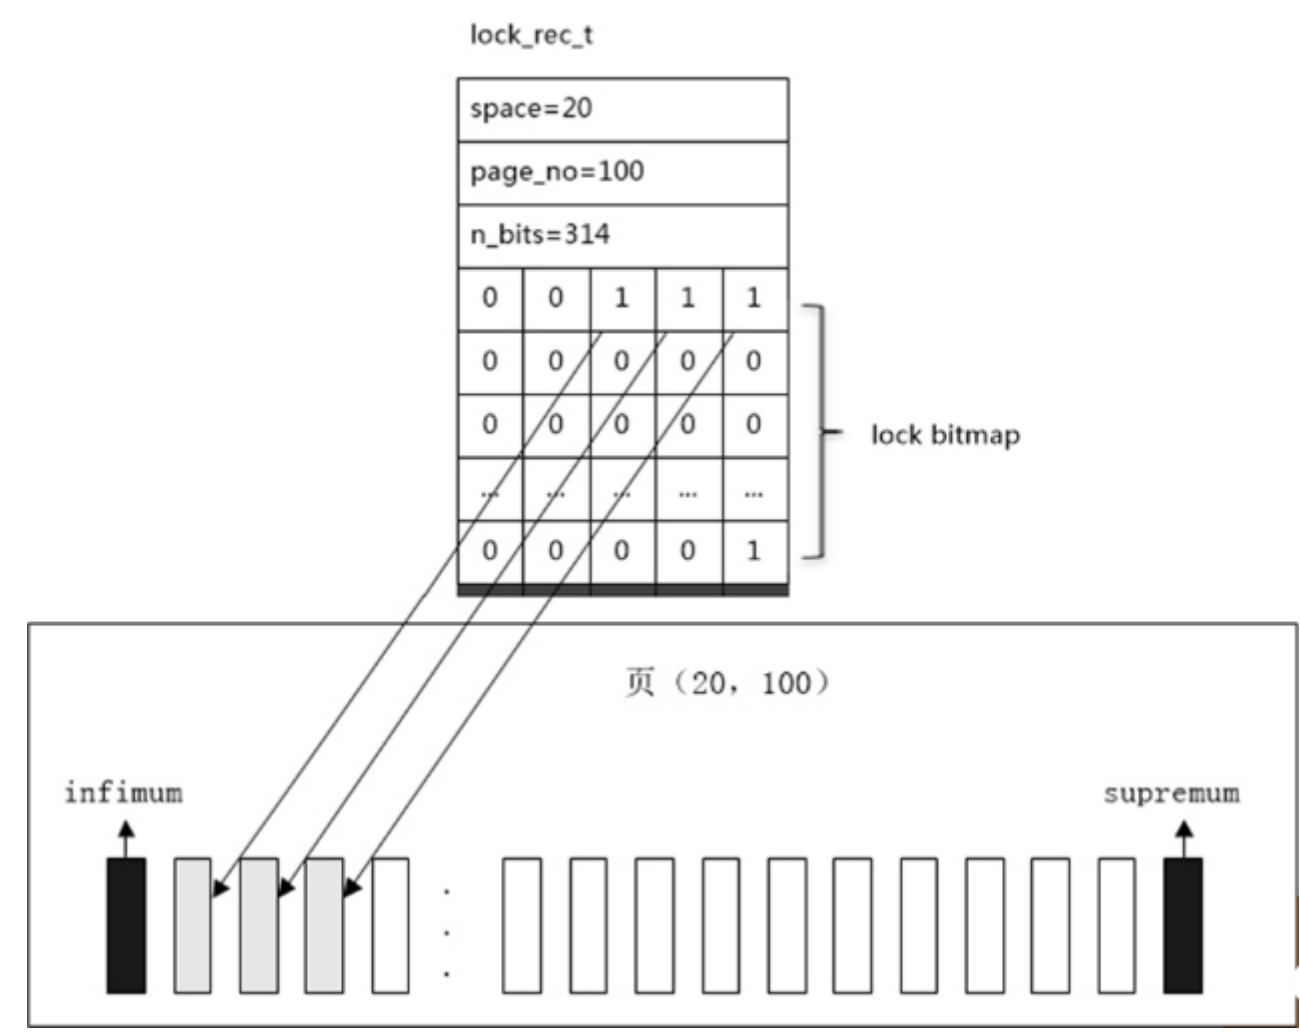
RECORD LOCKS space id 143 page no 3 n bits 72 index GEN_CLUST_INDEX of table 'burn_test'.'t_lock_3' trx id 26223 lock_mode X locks rec but not gap -- index GEN_CLUST_INDEX 聚集索引上的 Record-Lock
Record lock, heap no 4 PHYSICAL RECORD: n_fields 4; compact format; info bits 0
0: len 6; hex 00000001206; asc  ;;
1: len 6; hex 00000006250; asc  bP;;
2: len 7; hex c5000002a012c; asc  * ,;;
3: len 4; hex 8000000d; asc  ;;

RECORD LOCKS space id 143 page no 4 n bits 72 index idx_a of table 'burn_test'.'t_lock_3' trx id 26223 lock_mode X locks gap before rec -- index idx_a 二级索引上的 Gap Lock（针对下一个记录20）
Record lock, heap no 5 PHYSICAL RECORD: n_fields 2; compact format; info bits 0
0: len 4; hex 80000014; asc  ;;
1: len 6; hex 00000001207; asc  ;;
```

RR与RC性能对比

### 三. 锁的算法（二）

- 每个事物每个页 一个锁对象
  - 约100个字节
- 通过位图存放锁信息
  - 内存占用少
- 没有锁升级



- 锁重用

```
--
-- 终端会话1
--

mysql> begin;
Query OK, 0 rows affected (0.00 sec)

mysql> select * from t_lock_1 where a=13 for update;
+-----+
| a |
+-----+
| 13 |
+-----+
1 row in set (0.00 sec)

mysql> select * from t_lock_1 where a=13 lock in share mode;
+-----+
| a |
+-----+
| 13 |
+-----+
1 row in set (0.00 sec)

--
-- 终端会话2
--

mysql> show engine innodb status\G
-----省略部分输出-----
2 lock struct(s), heap size 1136, 1 row lock(s)
MySQL thread id 19, OS thread handle 140333220423424, query id 745 localhost root cleaning up
TABLE LOCK table 'burn_test'.'t_lock_1' trx id 26224 lock_mode IX
RECORD LOCKS space id 140 page no 3 n bits 72 index PRIMARY of table 'burn_test'.'t_lock_1' trx id 26224 lock_mode X locks rec but not gap -- 只有一把锁，不会出现两个锁，进一步降低了开销
Record lock, heap no 4 PHYSICAL RECORD: n_fields 3; compact format; info bits 0
0: len 4; hex 8000000d; asc  ;;
1: len 6; hex 0000000602b; asc  '+';;
2: len 7; hex ab000000470128; asc  G (;;
```

- 隐式锁

```
--

-- 终端会话1
--

mysql> begin;
Query OK, 0 rows affected (0.00 sec)

mysql> insert into t_lock_1 values(12);
Query OK, 1 row affected (0.00 sec)
-- 此时未提交，记录12上面应该有锁

--

-- 终端会话2
--

mysql> show engine innodb status\G -- 但是这里只能看到在表上的IX锁，而看不到记录12的记录锁
-- 因为show出来的都是显示的锁（已经创建了内存对象的）
-- 此时记录12中持有的是一把 隐式锁

-- -----省略部分输出-----
1 lock struct(s), heap size 1136, 0 row lock(s), undo log entries 1
MySQL thread id 19, OS thread handle 140333220423424, query id 750 localhost root cleaning up
TABLE LOCK table 'burn_test'.'t_lock_1' trx id 26225 lock_mode IX

--

-- 终端会话3
--

mysql> select * from t_lock_1 where a=12 for update; -- 另外一个事物中去修改a=12的这条未提交的记录
-- waiting.....

-- 此时 隐式锁 转换成了 显示锁

--

-- 终端会话2
--

mysql> show engine innodb status\G -- 此时就显示了锁的信息
-- -----省略部分输出-----
mysql tables in use 1, locked 1
LOCK WAIT 2 lock struct(s), heap size 1136, 2 row lock(s)
MySQL thread id 20, OS thread handle 140333220157184, query id 801 localhost root statistics
select * from t_lock_1 where a=12 for update
----- TRX HAS BEEN WAITING 2 SEC FOR THIS LOCK TO BE GRANTED:
RECORD LOCKS space id 140 page no 3 n bits 72 index PRIMARY of table 'burn_test'.'t_lock_1' trx id 26226 lock_mode X locks rec but not gap waiting -- 新的事物在等待这个 Record Lock
Record lock, heap no 6 PHYSICAL RECORD: n_fields 3; compact format; info bits 0
0: len 4; hex 8000000c; asc ;;
1: len 6; hex 000000006671; asc fq;;
2: len 7; hex df000000270110; asc ' ;;

-----
TABLE LOCK table 'burn_test'.'t_lock_1' trx id 26226 lock_mode IX
RECORD LOCKS space id 140 page no 3 n bits 72 index PRIMARY of table 'burn_test'.'t_lock_1' trx id 26226 lock_mode X locks rec but not gap waiting
Record lock, heap no 6 PHYSICAL RECORD: n_fields 3; compact format; info bits 0
0: len 4; hex 8000000c; asc ;;
1: len 6; hex 000000006671; asc fq;;
2: len 7; hex df000000270110; asc ' ;;

---TRANSACTION 26225, ACTIVE 313 sec
2 lock struct(s), heap size 1136, 1 row lock(s), undo log entries 1
MySQL thread id 19, OS thread handle 140333220423424, query id 750 localhost root cleaning up
TABLE LOCK table 'burn_test'.'t_lock_1' trx id 26225 lock_mode IX
RECORD LOCKS space id 140 page no 3 n bits 72 index PRIMARY of table 'burn_test'.'t_lock_1' trx id 26225 lock_mode X locks rec but not gap -- Record lock
Record lock, heap no 6 PHYSICAL RECORD: n_fields 3; compact format; info bits 0
0: len 4; hex 8000000c; asc ;;
1: len 6; hex 000000006671; asc fq;;
2: len 7; hex df000000270110; asc ' ;;

-- 尽管此时还没有提交（没有创建内存对象），但是可以通过 trx_id 来判断该记录上是否有锁；
-- 通过 在线事物列表 中是否存在该 trx_id 进行判断，因为事物没有提交（即活跃），就表示在这个记录上是有锁的
```

四. 一致性的非锁定读

在隔离级别为 **RU**、**RC**、**RR** 时，读取（select）操作是不会加锁的，通过行多版本控制（**MVCC**）的方式读取当前执行时间点的记录。

```
--

-- 终端会话1
--

mysql> begin;
Query OK, 0 rows affected (0.00 sec)

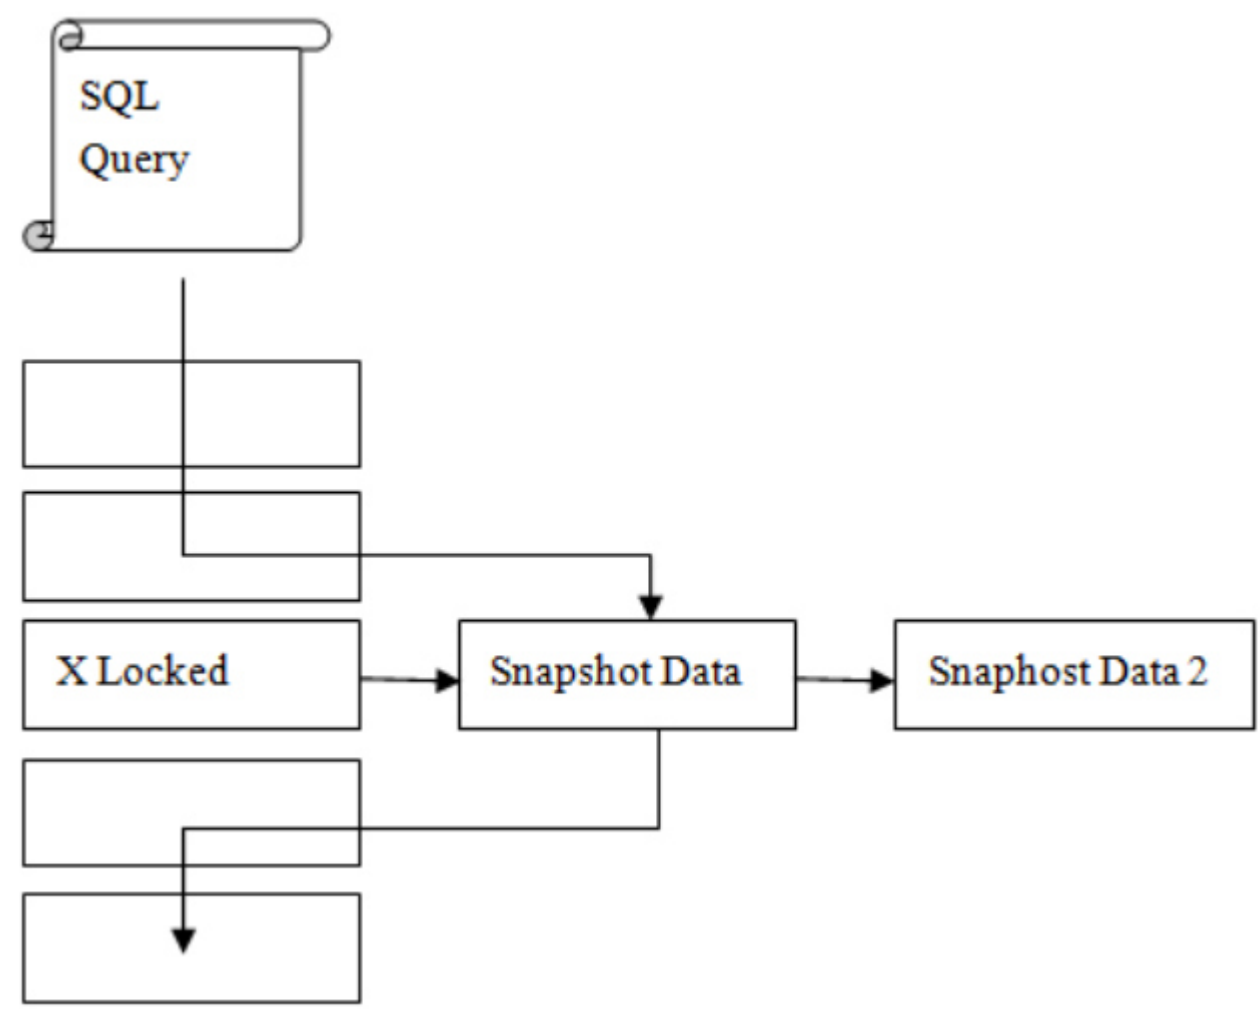
mysql> select * from t_lock_1 where a=13; -- 非锁定的读，不会阻塞
+-----+
| a |
+-----+
| 13 |
+-----+
1 row in set (0.00 sec)

--

-- 终端会话2
--

mysql> begin;
Query OK, 0 rows affected (0.00 sec)

mysql> select * from t_lock_1 where a=13; -- 非锁定的读，不会阻塞
+-----+
| a |
+-----+
| 13 |
+-----+
1 row in set (0.00 sec)
```



通过 **UNDO** 指针的指向，可以读取前一个版本甚至前几个版本的记录（*即通过UNDO来构造版本记录*），从而实现 **快照读**（**Snapshot Read**）

通过 **trx\_id** 判断该记录是否被锁住（在线事物列表），从而决定是否要读取之前的版本（**UNDO**）

五. 死锁

5.1. 死锁的介绍

- 两个或两个以上的事物在执行过程中，因争夺锁资源而造成的一种互相等待的现象
  - AB-BA、A和B互相等待
- 解决死锁
  - 超时，**死锁**和**锁超时**不是同一个东西（*锁超时是解决死锁的一种方式*）
    - --innodb\_lock\_timeout，默认是50s
- wait-for graph
  - 自动死锁检测

数据库中的死锁和程序中的死锁不同，数据库中的死锁是不能完全避免的，且数据库中的死锁有检测机制

5.2. AB-BA死锁演示

```
--

-- 终端会话1
--

mysql> begin; -- step_1
Query OK, 0 rows affected (0.00 sec)

mysql> delete from t_lock_1 where a=10; -- step_3
Query OK, 1 row affected (0.00 sec)

mysql> delete from t_lock_1 where a=11; -- step_6
ERROR 1213 (40001): Deadlock found when trying to get lock; try restarting transaction -- 出现死锁

--

-- 终端会话2
--

mysql> begin; -- step_2
Query OK, 0 rows affected (0.00 sec)

mysql> delete from t_lock_1 where a=11; -- step_4
Query OK, 1 row affected (0.00 sec)

mysql> delete from t_lock_1 where a=10; -- step_5
Query OK, 1 row affected (3.98 sec)
```



在MySQL出现错误（非死锁）的时候，线程中的事物是中止的，需要显示提交或回滚；而死锁时，MySQL会自动回滚

假如一个事物中有10个步骤，当执行到步骤8时

1. 出现普通错误，在步骤8中停止，只要重新执行步骤8即可，然后再继续执行后续操作
2. 出现死锁时，该事物中，步骤8之前的操作都会自动回滚

```
mysql> show engine innodb status\G      -- 显示最近的死锁信息，查看后就看不到了
-- -----省略部分输出-----
LATEST DETECTED DEADLOCK
-----
2016-02-13 15:23:27 0x7fa1dfbcd700
*** (1) TRANSACTION:
TRANSACTION 26240, ACTIVE 9 sec starting index read
mysql tables in use 1, locked 1
LOCK WAIT 3 lock struct(s), heap size 1136, 2 row lock(s), undo log entries 1
MySQL thread id 23, OS thread handle 140333220423424, query id 963 localhost root updating
delete from t_lock_1 where a=10
*** (1) WAITING FOR THIS LOCK TO BE GRANTED:
RECORD LOCKS space id 140 page no 3 n bits 72 index PRIMARY of table 'burn_test'.'t_lock_1' trx id 26240 lock_mode X locks rec but not gap waiting
Record lock, heap no 2 PHYSICAL RECORD: n_fields 3; compact format; info bits 32
0: len 4; hex 8000000a; asc      ;;
1: len 6; hex 00000000667f; asc      f ;;
2: len 7; hex 67000000330ab9; asc g   3 ;;

*** (2) TRANSACTION:
TRANSACTION 26239, ACTIVE 14 sec starting index read, thread declared inside InnoDB 5000
mysql tables in use 1, locked 1
3 lock struct(s), heap size 1136, 2 row lock(s), undo log entries 1
MySQL thread id 22, OS thread handle 140333220157194, query id 964 localhost root updating
delete from t_lock_1 where a=11
*** (2) HOLDS THE LOCK(S):
RECORD LOCKS space id 140 page no 3 n bits 72 index PRIMARY of table 'burn_test'.'t_lock_1' trx id 26239 lock_mode X locks rec but not gap
Record lock, heap no 2 PHYSICAL RECORD: n_fields 3; compact format; info bits 32
0: len 4; hex 8000000a; asc      ;;
1: len 6; hex 00000000667f; asc      f ;;
2: len 7; hex 67000000330ab9; asc g   3 ;;

*** (2) WAITING FOR THIS LOCK TO BE GRANTED:
RECORD LOCKS space id 140 page no 3 n bits 72 index PRIMARY of table 'burn_test'.'t_lock_1' trx id 26239 lock_mode X locks rec but not gap waiting
Record lock, heap no 3 PHYSICAL RECORD: n_fields 3; compact format; info bits 32
0: len 4; hex 8000000b; asc      ;;
1: len 6; hex 000000006680; asc      f ;;
2: len 7; hex 68000000330259; asc h   3 Y;;

*** WE ROLL BACK TRANSACTION (2)
-----
TRANSACTIONS
-----
Trx id counter 26246
Purge done for trx's n:o < 26245 undo n:o < 0 state: running but idle
History list length 316
LIST OF TRANSACTIONS FOR EACH SESSION:
---TRANSACTION 421809722470000, not started
0 lock struct(s), heap size 1136, 0 row lock(s)
---TRANSACTION 421809722468176, not started
0 lock struct(s), heap size 1136, 0 row lock(s)
---TRANSACTION 26240, ACTIVE 748 sec
3 lock struct(s), heap size 1136, 2 row lock(s), undo log entries 2
MySQL thread id 23, OS thread handle 140333220423424, query id 963 localhost root cleaning up
TABLE LOCK table 'burn_test'.'t_lock_1' trx id 26240 lock mode IX
RECORD LOCKS space id 140 page no 3 n bits 72 index PRIMARY of table 'burn_test'.'t_lock_1' trx id 26240 lock_mode X locks rec but not gap
Record lock, heap no 3 PHYSICAL RECORD: n_fields 3; compact format; info bits 32
0: len 4; hex 8000000b; asc      ;;
1: len 6; hex 000000006680; asc      f ;;
2: len 7; hex 68000000330259; asc h   3 Y;;

RECORD LOCKS space id 140 page no 3 n bits 72 index PRIMARY of table 'burn_test'.'t_lock_1' trx id 26240 lock_mode X locks rec but not gap
Record lock, heap no 2 PHYSICAL RECORD: n_fields 3; compact format; info bits 32
0: len 4; hex 8000000a; asc      ;;
1: len 6; hex 000000006680; asc      f ;;
2: len 7; hex 6800000033027b; asc h   3 ;;

mysql> show variables like 'innodb%dead%';
+-----+
| Variable_name | Value |
+-----+
| innodb_print_all_deadlocks | ON    | -- 会将死锁信息打印到err_log（搜索 deadlock 关键字）
+-----+
1 row in set (0.00 sec)
```